

# Lab 7

Xiaoyu Lian (A17025943)

Today we are going to learn how to apply different machine learning methods, beginning with clustering:

The goal here is to find groups/ clusters in your input data


first I will make up some data with clear groups. For this I will use the `rnorm()` function:

```
rnorm(10)
```

```
[1] -0.80874840  0.22000169 -1.00164155  0.30328705 -0.01129572  0.22424665  
[7] -1.49504265 -0.48315624 -0.18176392 -1.23023561
```


```
hist(rnorm(10000, mean = 3))
```

**Histogram of `rnorm(10000, mean = 3)`**



```
a <- hist(c(rnorm(10000,3),
            rnorm(10000,-3)))
```

## Histogram of c(rnorm(10000, 3), rnorm(10000, -3))



```
a
```

```
$breaks
```

```
[1] -7 -6 -5 -4 -3 -2 -1 0 1 2 3 4 5 6 7
```

```
$counts
```

```
[1] 16 211 1388 3385 3429 1351 229 228 1369 3455 3418 1306 196 19
```

```
$density
```

```
[1] 0.00080 0.01055 0.06940 0.16925 0.17145 0.06755 0.01145 0.01140 0.06845
[10] 0.17275 0.17090 0.06530 0.00980 0.00095
```

```
$mids
```

```
[1] -6.5 -5.5 -4.5 -3.5 -2.5 -1.5 -0.5 0.5 1.5 2.5 3.5 4.5 5.5 6.5
```

```
$xname
```

```
[1] "c(rnorm(10000, 3), rnorm(10000, -3))"
```

```
$equidist
```


```
[1] TRUE
```

```
attr("class")  
[1] "histogram"
```

```
n <- 30  
x <- c(rnorm(n,3), rnorm (n, -3))  
y<- rev(x)  
z<- cbind(x,y)  
head(z)
```

```
      x      y  
[1,] 2.964549 -2.530812  
[2,] 3.102970 -1.824010  
[3,] 3.924798 -1.714268  
[4,] 3.407469 -4.420817  
[5,] 2.310251 -3.061530  
[6,] 3.897858 -2.864947
```

```
hist(z)
```



Use the `kmeans()` function setting `k` to 2 and `nstart = 20` inspect the results `>Q`. How many points are in each cluster `>Q`. What 'component' of your result object details - cluster size - cluster assignment - cluster mist

plot `x` colored by the `kmeans` cluster assignment and add cluster centers as blue points

```
km <- kmeans(z, centers = 2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-2.815180	2.905294
2	2.905294	-2.815180

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 60.03475 60.03475
(between_SS / total_SS = 89.1 %)
```

Available components:

[1] "cluster"	"centers"	"totss"	"withinss"	"tot.withinss"
[6] "betweenss"	"size"	"iter"	"ifault"	

Results in `kmeans` object `km`

```
attributes(km)
```

\$names				
[1] "cluster"	"centers"	"totss"	"withinss"	"tot.withinss"
[6] "betweenss"	"size"	"iter"	"ifault"	

```
$class
[1] "kmeans"
```

cluster size

```
km$totss
```

```
[1] 1101.784
```

number of points in each cluster

```
km$size
```

```
[1] 30 30
```

cluster assignment / membership


```
km$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 1  
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Q. plot `z` colored by the `kmeans` cluster assignment and add cluster centers as blue points

It will recluster the shorter color vector to be the same length as the larger (numbers of data point in `z`)

```
plot(z, col = c("red", "blue"))
```




```
plot(z, col = km$cluster)
```




We can use the `points()` function to add new points to an existing plot...like the cluster centers

```
plot(z, col = km$cluster)
points(km$centers, col = "blue", pch = 15, cex = 3)
```



Q. Can you run kmeans and ask for 4 clusters please and plot the results like we have done above?

```
km2 <- kmeans(z, centers = 4)
plot(z, col = km2$cluster)
points(km2$centers, col = "blue", pch = 15)
```



## Hierarchical Clustering

Let's take our some made-up data **z** and see how cluster work

first we need a distance matrix of our data to be clustered

```
d <- dist(z)
hc <- hclust(d)
hc
```

Call:


```
hclust(d = d)
```

```
Cluster method   : complete
Distance         : euclidean
Number of objects: 60
```

```
plot(hc)
abline(h = 3, col = "red")
```



## Cluster Dendrogram




d  
hclust (\*, "complete")

I can get my cluster membership vector by “cutting the tree” with the `cutree()` function

```
grps <- cutree(hc, h = 8)
```

Can you plot `z` colored by our hclust results:

```
plot(z, col = grps)
```



Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions? Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
```


```
dim(x)
```

```
[1] 17  4
```

```
head(x)
```


	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```




Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



A so-called “Pairs” plot can be useful for small datasets like this one >Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



```
#Country name is the x-axis for all the graphs in that column
```

It is hard to see structure and trends in even this small data set. How will we ever do this when we have big data set with 1000s or 10s of thousands of things we are measuring...

###PCA to Rescue

Lets see how PCA deals with this dataset. So main function in base R to do PCA is called `prcomp()`

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Let's see what is inside this `pca` object that we created from running `prcomp()`

```
attributes(pca)
```

```
$names
```

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```


```
$class
```

```
[1] "prcomp"
```


```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
plot(pca$x[,1], pca$x[,2], col = c("black", "red", "blue", "green"), pch = 16,  
      xlab = "PC1", ylab = "PC2")
```



```
plot(pca$x[,1], pca$x[,2], col = c("black", "red", "blue", "green"),
     xlab = "PC1", ylab = "PC2")
text(pca$x[,1], pca$x[,2], colnames(x), col = c("black", "red", "blue", "green"))
```




```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```

```
[1] 67 29 4 0
```

```
z <- summary(pca)
z$importance
```

	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	2.921348e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00

```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



```
## Lets focus on PC1 as it accounts for > 90% of variance
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,1], las=2 )
```

